

**HCMUS - VNUHCM / FIT /**

**Computer Vision & Cognitive Cybernetics Department**

**Digital Image and Video Processing Application**

**Student ID:** 21127690

**Student name:** Ngo Nguyen Thanh Thanh

## **Report: Generative Adversarial Networks**

### **I. Evaluation summary:**

No	Task	Implementation	Completion (%)
1	Setup Google Colab for training	Configure Colab environment, install necessary libraries	100%
2	Load and preprocess dataset	Use MNIST dataset, normalize images	100%
3	Define Generator model	Implement a neural network to generate images	100%
4	Define Discriminator model	Implement a classifier to distinguish real vs fake images	100%
5	Setup optimizers for GAN training	Use Adam optimizer for both Generator and Discriminator	100%
6	Train GAN model	Train Generator and Discriminator iteratively	100%
7	Evaluate model performance	Create the loss curve analysis helps monitor GAN training stability	100%
8	Generate and visualize images	Implement function to generate and display images using trained Generator	100%

### **II. List of features and file structure:**

#### **Functions and Methods Used**

The notebook primarily utilizes the following libraries:

- torch, torch.nn, torch.optim: For defining and training neural networks.
- torchvision.transforms: For preprocessing image data.
- matplotlib.pyplot, numpy: For visualization and numerical operations.

#### **Main Functions**

The main functions in the notebook can be categorized as follows:

## 1. Model Implementation

This section includes the definitions of the **Discriminator (D)** and **Generator (G)** models.

- **Discriminator (D)**
  - `class Discriminator(nn.Module):` Defines a neural network for distinguishing real and generated images.
  - `__init__(self, inp_dim=784):` Initializes the discriminator with fully connected layers.
  - `forward(self, x):` Passes input through the network and outputs a probability score.
- **Generator (G)**
  - `class Generator(nn.Module):` Defines a neural network for generating synthetic images.
  - `__init__(self, z_dim=100):` Initializes the generator with a latent space input.
  - `forward(self, x):` Transforms random noise into an image representation.

## 2. Data Processing

Handles image dataset loading and preprocessing.

- `transforms.ToTensor():` Converts image data into tensors.
- `transforms.Normalize((0.5,), (0.5,)):` Normalizes the dataset for stable training.
- `x = x.view(x.size(0), 784):` Flattens the images for the neural network input.

## 3. Training Functions

Manages the training process of the GAN model.

- `optimizerD = torch.optim.Adam(D.parameters(), lr=0.0002):` Optimizer for the discriminator.
- `optimizerG = torch.optim.Adam(G.parameters(), lr=0.0002):` Optimizer for the generator.
- `lossD = criterion(output, label):` Computes the loss for the discriminator.
- `lossG = criterion(output, label):` Computes the loss for the generator.
- `lossD.backward(), lossG.backward():` Performs backpropagation.
- `optimizerD.step(), optimizerG.step():` Updates model weights.

## 4. Result Visualization

After training, the generated images are displayed.

- `make_grid(images, nrow=8, normalize=True):` Creates a grid of generated images.
- `plt.imshow(...):` Displays the generated images.

# How to Run the Code

### 1. Setup Environment

- The code is designed to run in **Google Colab**.
- Mount Google Drive using:

```
from google.colab import drive
drive.mount('/content/drive')
```

- Ensure CUDA is available for GPU acceleration:

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)
```

## 2. Execute the Notebook Cells

- Run all cells sequentially to:
  - Load dependencies.
  - Define and initialize the **Discriminator** and **Generator**.
  - Preprocess the dataset.
  - Train the model with **backpropagation and optimization**.
  - Visualize generated images.

## Image proof:

✓ 1. Cài đặt thư viện cần thiết

```
[39] import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torchvision.utils import make_grid
import matplotlib.pyplot as plt

import numpy as np
```

Double-click (or enter) to edit

✓ 2. Kiểm tra GPU

```
[40] device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)
```

Device: cuda

✓ 3. Cài đặt drive lưu kết quả và tải dataset

```
[41] from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[42] WORKING_DIR = '/content/drive/MyDrive/lab2-adip'
```

```
[43] import os

if not os.path.exists(WORKING_DIR):
    os.makedirs(WORKING_DIR)

%cd $WORKING_DIR
```

/content/drive/MyDrive/lab2-adip

## 4. Tải và lưu dataset MNIST

```
# Đường dẫn lưu dataset trên Google Drive
location_path = '/content/drive/MyDrive/' + 'lab2-adip/dataset'

# Định nghĩa transform trước khi truyền vào dataset
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# Batch size mới
batch_size = 128

# Tải dataset MNIST
dataset = torchvision.datasets.MNIST(root=location_path, train=True, download=True, transform=transform)

# Tạo DataLoader
dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size, shuffle=True)
```

## 5. Định nghĩa mô hình Discriminator

### Tự code

```
[45] class Discriminator(nn.Module):
    def __init__(self, inp_dim=784):
        super(Discriminator, self).__init__()
        self.w1 = nn.Parameter(torch.randn(inp_dim, 128) * 0.02) # Dùng nn.Parameter
        self.b1 = nn.Parameter(torch.zeros(128))
        self.w2 = nn.Parameter(torch.randn(128, 1) * 0.02)
        self.b2 = nn.Parameter(torch.zeros(1))

    def forward(self, x):
        x = x.view(x.size(0), 784)
        h = torch.matmul(x, self.w1) + self.b1
        h = torch.maximum(0.2 * h, h) # LeakyReLU
        out = torch.matmul(h, self.w2) + self.b2
        out = torch.sigmoid(out) # Sigmoid
        return out
```

### Tham khảo lab

```
class Discriminator(nn.Module):
    def __init__(self, inp_dim=784):
        super(Discriminator, self).__init__()
        self.fc1 = nn.Linear(inp_dim, 128)
        self.nonlin1 = nn.LeakyReLU(0.2)
        self.fc2 = nn.Linear(128, 1)

    def forward(self, x):
        x = x.view(x.size(0), 784) # Flatten (batch_size x 1 x 28 x 28) -> (batch_size x 784)
        h = self.nonlin1(self.fc1(x))
        out = self.fc2(h)
        out = torch.sigmoid(out)
        return out
```

## 6. Định nghĩa mô hình Generator

### ▼ Tự code

```
[46] class Generator(nn.Module):
    def __init__(self, z_dim=100):
        super(Generator, self).__init__()
        self.w1 = nn.Parameter(torch.randn(z_dim, 128) * 0.02)
        self.b1 = nn.Parameter(torch.zeros(128))
        self.w2 = nn.Parameter(torch.randn(128, 784) * 0.02)
        self.b2 = nn.Parameter(torch.zeros(784))

    def forward(self, x):
        h = torch.matmul(x, self.w1) + self.b1
        h = torch.maximum(0.2 * h, h) # LeakyReLU
        out = torch.matmul(h, self.w2) + self.b2
        out = torch.tanh(out) # [-1, 1]
        out = out.view(out.size(0), 1, 28, 28)
        return out
```

### ▼ Tham khảo lab

```
class Generator(nn.Module):
    def __init__(self, z_dim=100):
        super(Generator, self).__init__()
        self.fc1 = nn.Linear(z_dim, 128)
        self.nonlin1 = nn.LeakyReLU(0.2)
        self.fc2 = nn.Linear(128, 784)

    def forward(self, x):
        h = self.nonlin1(self.fc1(x))
        out = self.fc2(h)
        out = torch.tanh(out) # Đưa giá trị về khoảng [-1, 1]
        out = out.view(out.size(0), 1, 28, 28) # Chuyển về kích thước ảnh
        return out
```

### ▼ 7. Khởi tạo mô hình

```
[47] # Re-initialize D, G (Assuming Discriminator and Generator classes are defined)
D = Discriminator().to(device)
G = Generator().to(device)
```

## ▼ 8. Hàm mất mát và bộ tối ưu hóa

✓ [48]

```
print("Discriminator parameters:", sum(p.numel() for p in D.parameters() if p.requires_grad))
print("Generator parameters:", sum(p.numel() for p in G.parameters() if p.requires_grad))

# Set up the optimizers for Discriminator and Generator
# Adam is better than SGD for this task
#optimizerD = torch.optim.SGD(D.parameters(), lr=0.03)
#optimizerG = torch.optim.SGD(G.parameters(), lr=0.03)

# Uncomment to use Adam optimizer instead
optimizerD = torch.optim.Adam(D.parameters(), lr=0.0002)
optimizerG = torch.optim.Adam(G.parameters(), lr=0.0002)

# Define the loss function BCE (Binary Cross-Entropy)
criterion = nn.BCELoss()
batch_size = 128
```

Discriminator parameters: 100609  
Generator parameters: 114064

## 9. Huấn luyện mô hình

### ▼ Nhập giá trị noise vector

```
[49] noise_dim = int(input("Nhập giá trị noise_dim: "))
print(f"Noise dimension: {noise_dim}")
```

Nhập giá trị noise\_dim: 100  
Noise dimension: 100

```
import torch
import matplotlib.pyplot as plt

# Số epochs
epochs = 50

# Danh sách lưu loss
lossD_list = []
lossG_list = []

for epoch in range(epochs):
    lossD_epoch = 0
    lossG_epoch = 0
    num_batches = 0

    for i, data in enumerate(dataloader):
        x_real, _ = data
        x_real = x_real.to(device)
```



```
# Labels cho dữ liệu thật và giả
lab_real = torch.ones((x_real.size(0), 1), device=device)
lab_fake = torch.zeros((x_real.size(0), 1), device=device)

# ---- Training Discriminator ----
optimizerD.zero_grad()
lossD_real = criterion(D(x_real), lab_real)

z = torch.randn(x_real.size(0), noise_dim, device=device)
x_gen = G(z).detach()
lossD_fake = criterion(D(x_gen), lab_fake)

lossD = lossD_real + lossD_fake
lossD.backward()
optimizerD.step()

# ---- Training Generator ----
optimizerG.zero_grad()
z = torch.randn(x_real.size(0), noise_dim, device=device)
x_gen = G(z)
lossG = criterion(D(x_gen), lab_real) # Generator muốn đánh lừa D

lossG.backward()
optimizerG.step()

# Tính tổng loss trong epoch
lossD_epoch += lossD.item()
lossG_epoch += lossG.item()
num_batches += 1
```

```

optimizerG.zero_grad()
z = torch.randn(x_real.size(0), noise_dim, device=device)
x_gen = G(z)
lossG = criterion(D(x_gen), lab_real) # Generator muốn đánh lừa D

lossG.backward()
optimizerG.step()

# Tính tổng loss trong epoch
lossD_epoch += lossD.item()
lossG_epoch += lossG.item()
num_batches += 1

# Lưu loss trung bình vào danh sách
lossD_list.append(lossD_epoch / num_batches)
lossG_list.append(lossG_epoch / num_batches)

print(f"Epoch [{epoch+1}/{epochs}] - Loss D: {lossD_list[-1]:.4f}, Loss G: {lossG_list[-1]:.4f}")

# ---- Vẽ đồ thị ----
plt.figure(figsize=(8, 5))
plt.plot(range(1, epochs+1), lossD_list, label="Discriminator Loss", color="red")
plt.plot(range(1, epochs+1), lossG_list, label="Generator Loss", color="blue")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss of GAN")
plt.legend()
plt.show()

```

... Epoch [1/50] - Loss D: 1.1026, Loss G: 0.6367

## 10. Sinh ảnh và hiển thị

```

from PIL import Image

def show_generated_images(G, num_images=16, noise_dim=100, device="cpu", save_path=None):
    G.eval() # Đặt Generator ở chế độ đánh giá
    with torch.no_grad():
        z = torch.randn(num_images, noise_dim, device=device)
        fake_images = G(z).detach().cpu()

    grid = make_grid(fake_images, normalize=True, nrow=4)
    np_img = np.transpose(grid.numpy(), (1, 2, 0)) # Chuyển về định dạng ảnh

    plt.figure(figsize=(6, 6))
    plt.imshow(np_img)
    plt.axis("off")

    if save_path:
        image = Image.fromarray((np_img * 255).astype(np.uint8)) # Convert to image
        image.save(save_path)
        print(f"Ảnh đã lưu tại: {save_path}")

    plt.show()
    G.train() # Đặt lại Generator về chế độ training

# Gọi hàm và lưu ảnh vào Google Drive
save_path = "/content/drive/My Drive/lab2-adip/generated_image.png" # Thay đổi đường dẫn nếu cần
show_generated_images(G, device=device, save_path=save_path)

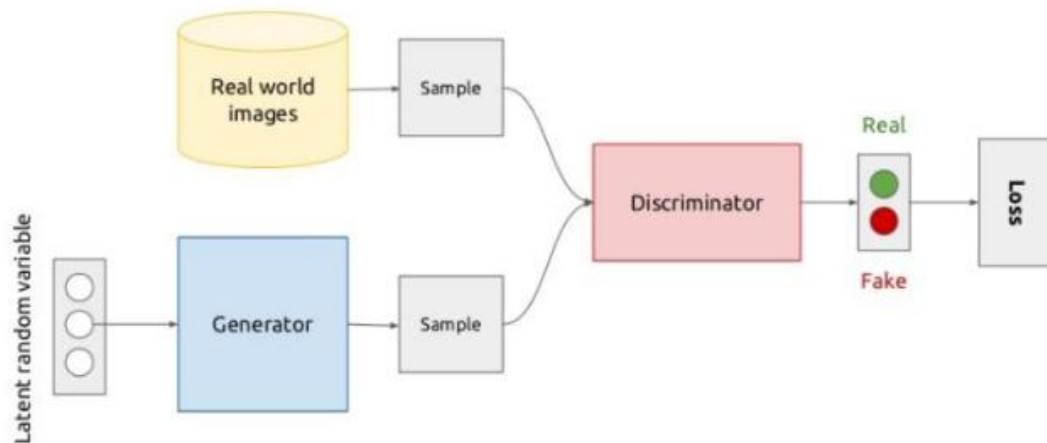
```

...

## III. Summarization of the usage



## Framework:



A generative adversarial network (GAN) uses two neural networks to compete with each other like in a game, one known as a “discriminator” and the other known as the “generator”.

- The Generator wants to learn to generate realistic images that are indistinguishable from the real data. The input of the Generator is a Gaussian noise random sample, and its output is a generated data point
- The Discriminator wants to tell the real & fake images apart. The input of the Discriminator is a datapoint or an image, and its output is a probability assigned to the datapoint being real. It can be seen as a binary classifier

## Detailed Usage and Algorithm Explanation:

### 1. Generative Adversarial Network (GAN) Overview

GAN consists of two models:

- **Discriminator  $D(x)$ :** Learns to classify real and fake images.
- **Generator  $G(z)$ :** Learns to generate realistic images from random noise  $z$ .

The objective function of GAN is:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

where:

- $p_{\text{data}}(x)$  is the real data distribution.
- $p_z(z)$  is the noise distribution used to generate fake samples.

## 2. Discriminator (D) Implementation and Explanation

### Mathematical Formulation

The discriminator is a binary classifier that takes an input  $x$  and outputs a probability  $D(x) \in [0,1]$  representing the likelihood that  $x$  is real. It is trained using **binary cross-entropy loss**:

$$\mathcal{L}_D = -\mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] - \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

### Pseudo Code for Discriminator Training

```
# Forward pass real images through Discriminator
real_output = D(real_images)
real_loss = criterion(real_output, torch.ones_like(real_output))

# Forward pass generated images through Discriminator
fake_images = G(noise)
fake_output = D(fake_images.detach())
fake_loss = criterion(fake_output, torch.zeros_like(fake_output))

# Compute total loss and backpropagate
lossD = real_loss + fake_loss
optimizerD.zero_grad()
lossD.backward()
optimizerD.step()
```

## 3. Generator (G) Implementation and Explanation

### Mathematical Formulation

The generator learns to transform random noise  $z$  into realistic images. It aims to maximize  $D(G(z))$  so that the discriminator classifies fake images as real:

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z(z)} [\log D(G(z))]$$

To improve stability, we use the alternative loss:

$$\mathcal{L}_G = \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

### Pseudo Code for Generator Training

```
# Generate fake images
fake_images = G(noise)
```

```

# Forward pass fake images through Discriminator
fake_output = D(fake_images)

lossG = criterion(fake_output, torch.ones_like(fake_output)) # Fool the discriminator

# Backpropagate
optimizerG.zero_grad()
lossG.backward()
optimizerG.step()

```

## 4. Training Process

### Overall Explanation

The training process follows the standard GAN approach, where we alternately update the **Discriminator (D)** and **Generator (G)** in each iteration.

- **Step 1:** Train **Discriminator (D)** to distinguish real and fake images.
- **Step 2:** Train **Generator (G)** to generate realistic images and fool the discriminator.

This min-max game continues until the generator produces high-quality images.

### Step 1: Training the Discriminator

**Goal:** The discriminator is trained to assign a high probability to real images and a low probability to generated images.

1. **Process real images**
  - Pass real images  $x$  through  $D(x)$
  - Compute loss using **binary cross-entropy (BCE)** with label  $y=1$  (real).
2. **Process fake images**
  - Generate fake images  $G(z)$  from random noise  $z$
  - Pass them through  $D(G(z))$
  - Compute loss using **BCE** with label  $y=0$  (fake).
3. **Update Discriminator**
  - Compute total loss  $\mathcal{L}_D$ .
  - Perform **backpropagation** and update  $D$  parameters.

### Pseudo Code for Discriminator Training

```

# Set Discriminator to training mode

```

```

D.train()

# Get real images from dataset
real_images, _ = next(iter(dataloader))

real_images = real_images.view(real_images.size(0), -1).to(device)

# Compute output for real images
real_output = D(real_images)

real_labels = torch.ones_like(real_output) # Real label = 1

real_loss = criterion(real_output, real_labels)

# Generate fake images
noise = torch.randn(batch_size, z_dim).to(device)

fake_images = G(noise).detach() # Stop gradient propagation to G

fake_output = D(fake_images)

fake_labels = torch.zeros_like(fake_output) # Fake label = 0

fake_loss = criterion(fake_output, fake_labels)

# Compute total Discriminator loss
lossD = real_loss + fake_loss

# Backpropagate and update D
optimizerD.zero_grad()

lossD.backward()

optimizerD.step()

```

## Step 2: Training the Generator

**Goal:** The generator learns to create realistic images so that the discriminator classifies them as real.

1. **Generate fake images**
  - Sample random noise  $z$  and pass it through  $G(z)$
2. **Trick the Discriminator**
  - Pass generated images through  $D(G(z))$
  - Instead of using label  $y=0$ , we use  $y=1$  (pretend fake images are real).
  - Compute loss using **BCE** with label  $y=1$  (fooling  $D$ ).
3. **Update Generator**
  - Compute total loss  $\mathcal{L}_G$ .
  - Perform **backpropagation** and update  $G$  parameters.

## Pseudo Code for Generator Training

```

# Set Generator to training mode
G.train()

# Generate fake images
noise = torch.randn(batch_size, z_dim).to(device)
fake_images = G(noise)

# Compute Discriminator's response to fake images
fake_output = D(fake_images)
fake_labels = torch.ones_like(fake_output) # Fool D into thinking fake images are real
lossG = criterion(fake_output, fake_labels)

# Backpropagate and update G
optimizerG.zero_grad()
lossG.backward()
optimizerG.step()

```

### Step 3: Complete Training Loop

**Goal:** Alternate between updating **D** and **G** for multiple epochs.

#### Full Training Loop Pseudo Code

```

for epoch in range(num_epochs):
    for real_images, _ in dataloader:
        # Train Discriminator
        real_images = real_images.view(real_images.size(0), -1).to(device)
        noise = torch.randn(batch_size, z_dim).to(device)
        fake_images = G(noise).detach()
        real_output = D(real_images)
        fake_output = D(fake_images)
        lossD_real = criterion(real_output, torch.ones_like(real_output))
        lossD_fake = criterion(fake_output, torch.zeros_like(fake_output))
        lossD = lossD_real + lossD_fake
        optimizerD.zero_grad()
        lossD.backward()
        optimizerD.step()

        # Train Generator
        noise = torch.randn(batch_size, z_dim).to(device)
        fake_images = G(noise)

```

```

fake_output = D(fake_images)

lossG = criterion(fake_output, torch.ones_like(fake_output))

optimizerG.zero_grad()

lossG.backward()

optimizerG.step()

# Print progress

print(f"Epoch [{epoch+1}/{num_epochs}], LossD: {lossD.item()}, LossG: {lossG.item()}")

```

## 5. Result Visualization

**Goal:** After training, visualize the **generated images** to assess the quality of the generator.

1. **Generate images** from random noise zzz.
2. **Convert tensor images** to a grid for visualization.
3. **Display the images** using matplotlib.

### Summary of the Training Process

1. **Training the Discriminator (D)**
  - Process real images and compute loss with  $y=1$
  - Generate fake images and compute loss with  $y=0$
  - Backpropagate and update D.
2. **Training the Generator (G)**
  - Generate fake images.
  - Compute loss by fooling D with  $y=1$
  - Backpropagate and update G
3. **Repeat the process** for multiple epochs until the generator produces realistic images.
4. **Visualize generated images** to evaluate the performance.

## Dataset: MNIST

### 1. Overview

The **MNIST (Modified National Institute of Standards and Technology) dataset** is a widely used benchmark dataset for handwritten digit recognition. It consists of **70,000 grayscale images** of handwritten digits from **0 to 9**, where:

- **60,000 images** are used for training.
- **10,000 images** are used for testing.

Each image has a resolution of **28 × 28 pixels** and is stored in **grayscale (1 channel)**. The pixel values range from **0 (black) to 255 (white)**.

### 2. Dataset Structure

- **Image Size:**  $28 \times 28 \times 28$  pixels
- **Number of Classes:** 10 (digits 0-9)
- **Color Mode:** Grayscale
- **Pixel Intensity Range:** 0-255 (normalized to 0-1 or -1 to 1 during preprocessing)
- **Data Split:**
  - **Training set:** 60,000 images
  - **Test set:** 10,000 images

### 3. Dataset Preprocessing


Before training, the images are preprocessed to **normalize** their pixel values and reshape them for input into the neural network.

#### Preprocessing steps:

1. Convert pixel values from **0-255 to 0-1** (or -1 to 1 for better stability).
2. Flatten each image into a **1D vector of 784 features** ( $28 \times 28 \times 28$ ).
3. Convert labels into tensors.

## IV. EXPERIMENTS AND EVALUATION:

**Epochs = 50, noise dim = 100:**



**Observations:**

1. **Discriminator Loss (D Loss):** Fluctuates between **0.6 and 1.2**, without a clear downward trend. This suggests that the discriminator is neither overpowering the generator nor converging smoothly.
2. **Generator Loss (G Loss):** Ranges from **1.0 to 1.6**, with noticeable oscillations. A high generator loss may indicate that it struggles to produce realistic samples.
3. **Training Dynamics:**
  - In the early epochs, the generator improves as the discriminator loss decreases.
  - Around the mid-training phase, instability increases, and neither model dominates.
  - By epoch 50, both losses remain erratic, which may suggest training instability or insufficient convergence.

#### Potential Issues & Fixes:


- **Unstable Training:** Loss fluctuations suggest the training process is not stable. Try **feature matching** or **label smoothing** to stabilize updates.
- **Mode Collapse:** If the generator produces limited variations, check generated images and consider techniques like **mini-batch discrimination**.
- **Hyperparameter Tuning:** Experiment with different **learning rates** (e.g., lowering generator LR) and **batch sizes** to improve stability.

#### Results:






Ảnh đã lưu tại: [/content/drive/My Drive/lab2-adip/generated\\_image.png](#)



Ảnh đã lưu tại: /content/drive/My Drive/lab2-adip/generated\_image.png



## Analysis of Generated Images

### Observations:

#### 1. Blurry and Noisy Outputs:

- The digits are recognizable but very **blurry and grainy**, indicating that the generator is struggling to produce sharp images.
- The background contains excessive **white noise**, suggesting unstable training or poor discriminator feedback.

#### 2. Mode Collapse Signs:

- Some digits appear **similar**, hinting that the generator may have collapsed to producing only a few variations.
- Some numbers are incomplete or distorted (e.g., the digit "1" in the middle row is overly thin).


#### 3. Training Instability:

- The low quality of details suggests **generator-discriminator imbalance**.
- The generator may not be learning meaningful features from the dataset.

### Potential Fixes:

- **Use Batch Normalization or Spectral Normalization** to stabilize training.
- **Adjust Learning Rates:** Reduce the generator's LR to improve fine details
- **Use Feature Matching Loss:** Helps the generator learn better structures
- **Train Longer:** The model might need more epochs for convergence.

**Epochs = 100, noise dim = 100**




## Observations


### 1. Training Loss Analysis:

- **Discriminator Loss (D Loss):** Fluctuates between 0.8 and 1.2, showing no clear downward trend. This suggests that the discriminator is neither overpowering the generator nor stabilizing effectively.
- **Generator Loss (G Loss):** Ranges from 1.4 to 2.0, with significant oscillations. The high loss indicates that the generator struggles to produce realistic samples.
- **Training Dynamics:**
  - In the early epochs, both losses fluctuate, indicating initial learning.
  - Around the mid-training phase, instability increases, and neither model dominates.
  - By epoch 100, both losses remain highly erratic, suggesting unstable training or insufficient convergence.

### Potential Issues & Fixes:

- **Unstable Training:** Loss fluctuations indicate instability. Try **feature matching** or **label smoothing** to stabilize updates.
- **Mode Collapse:** If the generator produces limited variations, check for repetition in generated images and consider **mini-batch discrimination**.
- **Hyperparameter Tuning:**
  - Reduce **learning rate** for the generator to prevent oscillations.
  - Experiment with **batch size** to balance updates.





## Analysis of Generated Images

### 1. Blurry and Noisy Outputs:

- The digits are somewhat recognizable but appear **blurry and grainy**, indicating that the generator struggles to refine details.
- The background contains excessive **white noise**, suggesting unstable training or poor discriminator feedback.

### 2. Signs of Mode Collapse:

- Some digits appear **similar**, meaning the generator might be producing only a few variations.
- Several numbers are **incomplete or distorted** (e.g., digit "1" in the middle row looks overly thin).

### 3. Training Instability:

- The low quality of generated details suggests a **generator-discriminator imbalance**.
- The generator may not be **learning meaningful features** from MNIST.

### Potential Fixes:

- **Use Batch Normalization or Spectral Normalization** to stabilize training.
- **Adjust Learning Rates:** Reduce the **generator's learning rate** to improve fine details.
- **Use Feature Matching Loss:** Helps the generator learn better structures.

- **Train Longer:** The model might need **more epochs** to reach better convergence.

## Comparison Between Custom Code and Lab Code & Its Impact on Results

### 1. Comparison of Discriminator Architecture

Component	Custom Code	Lab Code
Parameter Initialization	Uses nn.Parameter with torch.randn	Uses nn.Linear
Parameter Management	Manually manages weight matrices and biases	nn.Linear automatically manages them
Hidden Activation	torch.maximium(0.2 * h, h) (manual LeakyReLU)	nn.LeakyReLU(0.2) (built-in)
Number of Layers	2 layers (w1, w2)	2 layers (fc1, fc2)
Output Activation	torch.sigmoid	torch.sigmoid

#### Impact on Performance

- **Custom Code:** Since nn.Parameter is manually declared, weight updates and backpropagation may not be as optimized as nn.Linear, which can affect training stability.
- **Lab Code:** nn.Linear automatically manages parameters, reducing errors due to incorrect tensor size calculations or manual updates.
- **Efficiency:** Using nn.Linear is computationally more efficient due to PyTorch's internal optimizations, leading to smoother training.

### 2. Comparison of Generator Architecture

Component	Custom Code	Lab Code
Parameter Initialization	Uses nn.Parameter with torch.randn	Uses nn.Linear
Parameter Management	Manually manages weight matrices and biases	nn.Linear automatically manages them
Hidden Activation	torch.maximium(0.2 * h, h) (manual LeakyReLU)	nn.LeakyReLU(0.2) (built-in)
Number of Layers	2 layers (w1, w2)	2 layers (fc1, fc2)
Output Activation	torch.tanh	torch.tanh

#### Impact on Performance

- **Custom Code:**
  - Manually managing parameters with nn.Parameter can make weight updates less optimized, potentially causing unstable training.

- The use of `torch.matmul(x, self.w1) + self.b1` instead of `nn.Linear` might slow down training due to lack of PyTorch's internal optimizations.
- **Lab Code:**
  - `nn.Linear` ensures better computational efficiency, leading to smoother forward and backward propagation.
  - Easier to maintain and less prone to errors.
- **Efficiency:** Like the Discriminator, using `nn.Linear` improves computational efficiency and ensures optimal training performance.



### 3. Overall Impact on Training and Performance

Factor	Custom Code	Lab Code
Training Speed	Slower due to manual parameter management	Faster due to <code>nn.Linear</code> optimizations
Training Stability	May be less stable due to manual updates	More stable due to PyTorch's built-in mechanisms
Readability & Maintainability	More difficult due to manual parameter handling	Easier to maintain with <code>nn.Linear</code>
Flexibility & Generalization	More error-prone when modifying architecture	More adaptable to changes

### Conclusion

- **Lab Code is more efficient in terms of speed, stability, and maintainability.**
- **Custom Code is useful for understanding GAN operations at a lower level but is harder to optimize and prone to errors.**
- **Using `nn.Linear` leads to faster and more stable training.**

## Results:





## Analysis of GAN Training Loss

### 1. Observations

- Both generator and discriminator losses fluctuate significantly, showing instability.
- No clear convergence; sharp spikes suggest training imbalance.

### 2. Potential Issues

- **Mode Collapse:** The generator may be producing limited outputs, causing loss oscillations.
- **Unstable Training:** High variance in updates, possibly due to learning rate or batch size settings.